

## B2: Build an Artificial Neural Network by implementing the Backpropagation algorithm and test the same using appropriate data sets.

```
In [8]: # import libraries
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score
```

```
In [9]: # Load the Iris dataset
iris = datasets.load_iris()
X = iris.data
y = iris.target
```

```
In [10]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```



```
In [12]: # Create and train the MLP model
# The 'adam' solver uses stochastic gradient descent + backpropagation to update weights.
mlp = MLPClassifier(hidden_layer_sizes=(10,), max_iter=1000, activation='relu', solver='adam', random_state=42)
```

Parameter	Meaning
hidden_layer_sizes=(10,)	One hidden layer with <b>10 neurons</b> . You can have multiple layers like (10, 5) for two layers: 10 neurons in first, 5 in second.
max_iter=1000	Train the model for a <b>maximum of 1000 iterations</b> (epochs) if convergence isn't reached earlier.
activation='relu'	Use the <b>ReLU (Rectified Linear Unit)</b> activation function: $f(x) = \max(0, x)$ . ReLU helps avoid vanishing gradients.
solver='adam'	Optimization algorithm used: <b>Adam</b> optimizer (Adaptive Moment Estimation). It combines RMSProp and SGD with momentum.
random_state=42	Ensures reproducibility by setting a <b>fixed seed</b> for random number generation.

In [13]: `mlp.fit(X_train, y_train)`

```
c:\Users\Venkat\AppData\Local\Continuum\anaconda3\envs\myenv_upgrade\Lib\site-packages\sklearn\netw
ork\_multilayer_perceptron.py:780: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1000) rea
ched and the optimization hasn't converged yet.
  warnings.warn(
```

Out[13]:

▼ MLPClassifier    
Parameters ([https://  
scikit-  
learn.org/1.7/  
modules/  
generated/](https://scikit-learn.org/1.7/modules/generated/)

In [14]: *#Make predictions*  
`y_pred = mlp.predict(X_test)`

In [15]: *#Evaluate the model*  
`accuracy = accuracy_score(y_test, y_pred)`  
`print(f"Accuracy: {accuracy * 100:.2f}%")`

Accuracy: 97.78%